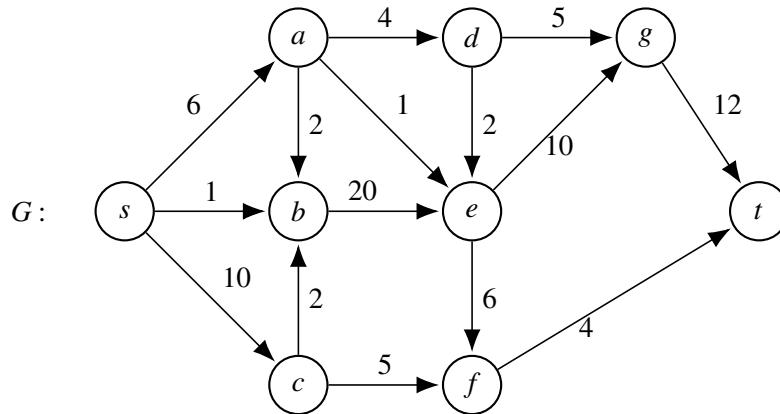


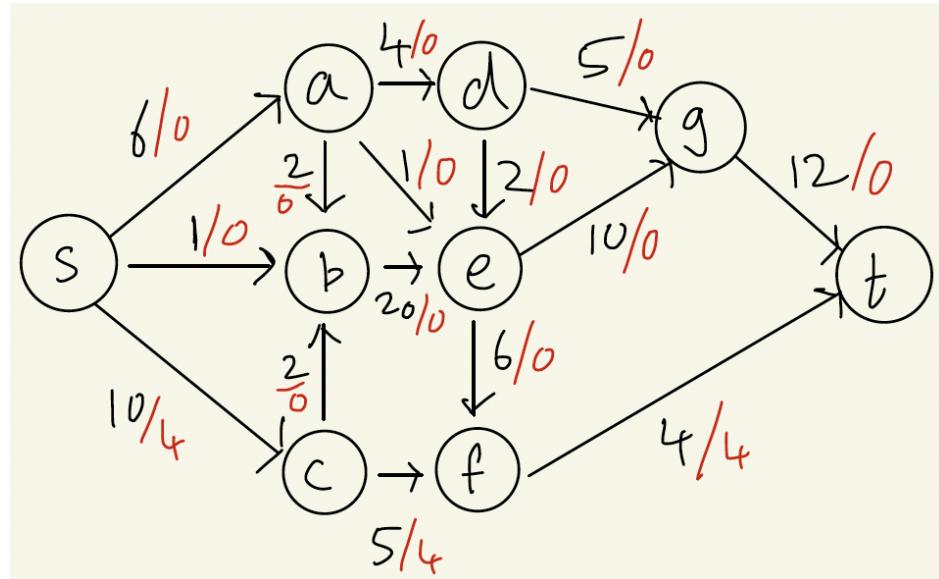
1. ( pts.) **Ford-Fulkerson Algorithm.** Use the Edmonds-Karp algorithm (namely, Ford-Fulkerson where each augmenting path is found by BFS) to find the maximum flow and the corresponding minimum cut in the given  $s - t$  flow network. Process neighbors in alphabetical order during BFS. Show the augmenting path and draw the residual graph  $G_f$  for each step. Show the final maximum-flow and the minimum cut.



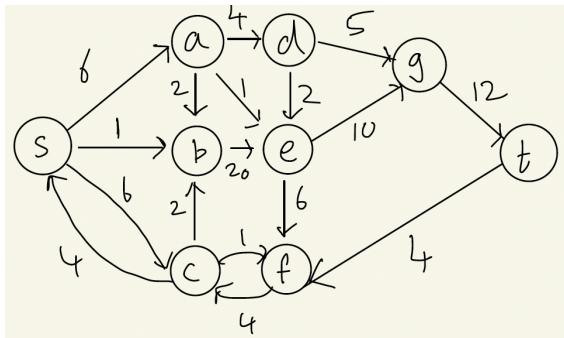
**Answer:**

The maximum flow is 13 units and the minimum  $(S, T)$ -cut is  $(\{S, C, F\}, \{A, B, D, E, G, T\})$ .

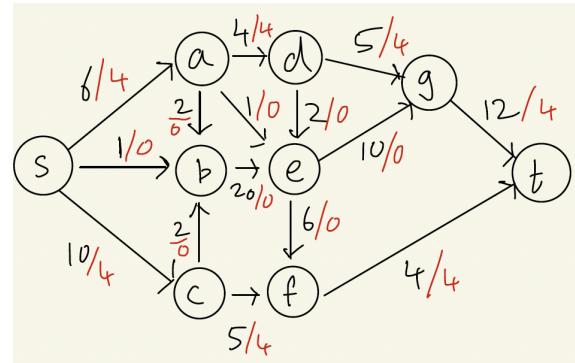
Using the Edmonds-Karp algorithm, the first  $s - t$  path chosen for augmenting the flow is  $s \rightarrow c \rightarrow f \rightarrow t$ . The bottleneck edge is  $f - t$  with a capacity of 4, so we augment the flow along this path by adding 4 units of flow to each edge. In the following figure, the flow values are written in red.



The new residual graph is shown below. Here, the  $s - t$  path chosen is  $s \rightarrow a \rightarrow d \rightarrow g \rightarrow t$ . The bottleneck edge is  $a - d$  with a capacity of 4, so we augment the flow along this path by adding 4 units.

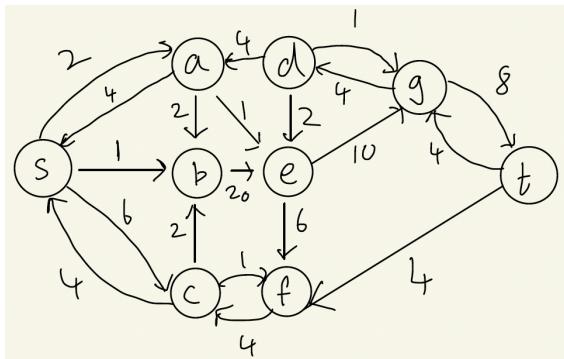


(a) Residual Graph

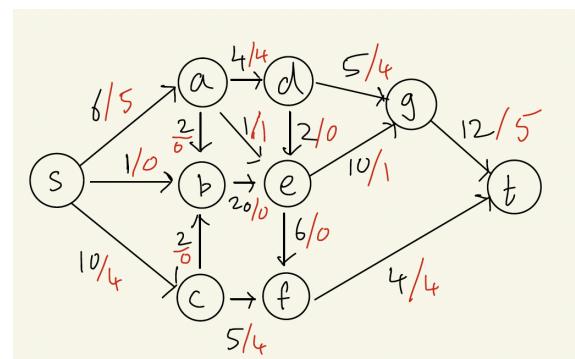


(b) Augmented Flow

After augmenting the flow, we get the following residual graph. In this iteration, the  $s - t$  path chosen is  $s \rightarrow a \rightarrow e \rightarrow g \rightarrow t$ . The bottleneck edge is  $a - e$  with a capacity of 1. So, we augment the flow along this path by adding 1 unit of flow.

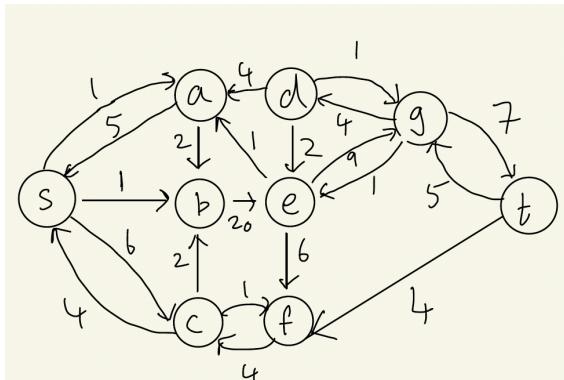


(a) Residual Graph

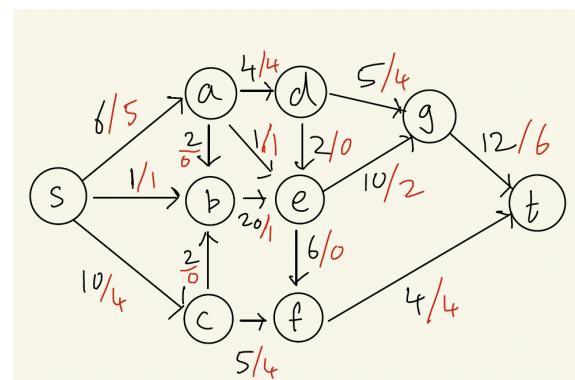


(b) Augmented Flow

The new residual graph is shown below. The  $s - t$  path chosen is  $s \rightarrow b \rightarrow e \rightarrow g \rightarrow t$ . The bottleneck edge is  $s - b$  with a capacity of 1. We augment the flow along this path by adding 1 unit of flow.

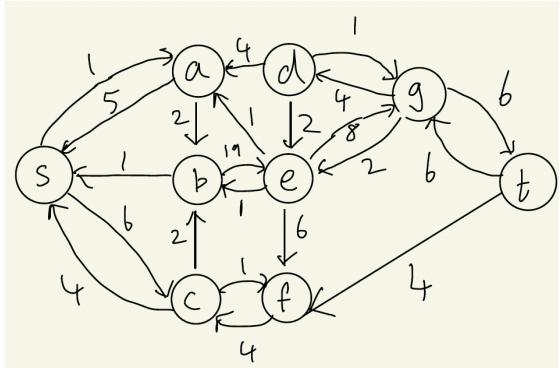


(a) Residual Graph

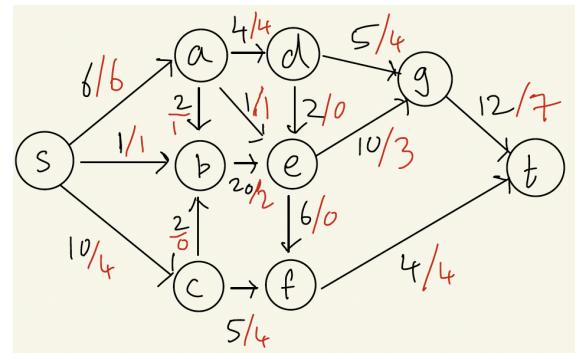


(b) Augmented Flow

The new residual graph is shown below. The  $s - t$  path chosen is  $s \rightarrow a \rightarrow b \rightarrow g \rightarrow e \rightarrow t$ . The bottleneck edge is  $s - a$  with a capacity of 1. We augment the flow along this path by adding 1 unit of flow.

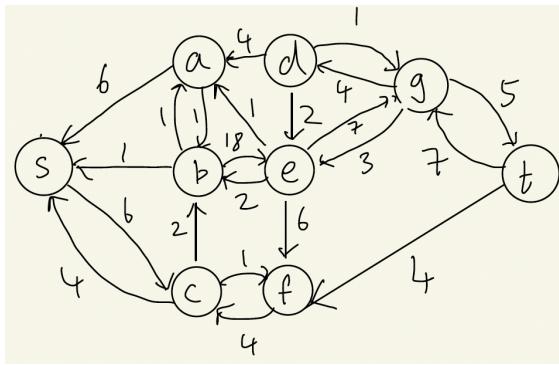


(a) Residual Graph

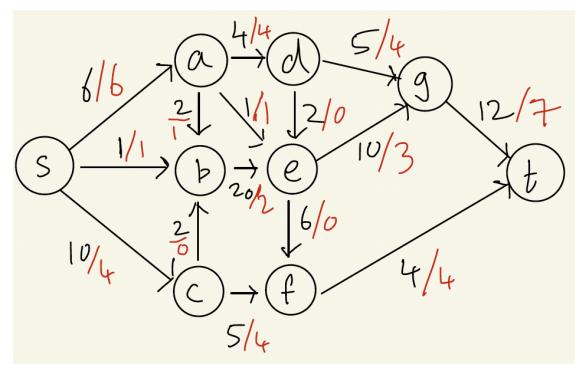


(b) Augmented Flow

The new residual graph is shown below. The  $s - t$  path chosen is  $s \rightarrow c \rightarrow b \rightarrow e \rightarrow g \rightarrow t$ . The bottleneck edge is  $c - b$  with a capacity of 2. We augment the flow along this path by adding 2 units of flow.



(a) Residual Graph

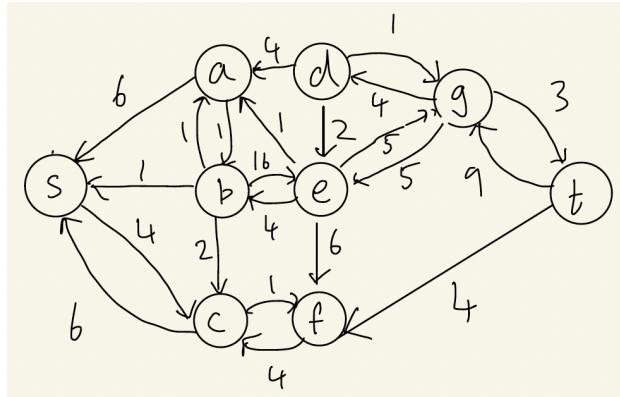


(b) Augmented Flow

The new residual graph is shown below. There is no  $s - t$  path so the algorithm terminates. The maximum flow is 13. By definition, the minimum cut is  $(S^*, T^*)$  where:

$$S^* = \{v \mid v \text{ is reachable from } s \text{ in } G_f^*\}$$

and  $T^* = V \setminus S^*$ .



- 2. ( pts.) Max Flow Formulation.** Consider a variation of the standard max-flow problem in a flow network. In addition to the existing edge capacities, each vertex has a capacity indicating the maximum amount of flow that is allowed to pass through. Show that this problem can be reduced to the standard max-flow

problem and explain any pre-processing or additional data structures used. [Hint: transform this graph with both edge and vertex capacities into a regular flow network with only edge capacities by doubling the number of vertices.]

**Answer:**

Construct a new network  $G'$  by splitting each vertex  $v$  of  $G$  into two vertices:  $v_{\text{in}}$  and  $v_{\text{out}}$ . Make all edges going into  $v$  in  $G$  go into  $v_{\text{in}}$  and all edges leaving  $v$  in  $G$  leave  $v_{\text{out}}$  in  $G'$ . Finally, insert an edge from  $v_{\text{in}}$  to  $v_{\text{out}}$  with capacity equal to the capacity of  $v$  in  $G$ . Now it is sufficient to solve the maximum flow problem with edge capacities in  $G'$ . This works as every valid flow through  $v$  in  $G$  that satisfies  $v$ 's capacity constraint can be made into a valid flow through  $v_{\text{in}}$  and  $v_{\text{out}}$  in  $G'$ , and vice versa.

3. ( pts.) **Critical Edge.** An edge of an  $s - t$  flow network is called *critical* if decreasing the capacity of this edge results in a decrease in the  $s - t$  maximum flow. Give an efficient algorithm that finds a critical edge in a network. Explain the correctness of your algorithm and analyze the running time.

**Answer:**

Let  $f > 0$  be the value of the maximum  $s - t$  flow (if  $f = 0$ , then there are no critical edges). By the max-flow min-cut theorem the capacity of any minimum capacity  $s - t$  cut is also  $f$ . Fix a minimum  $s - t$  cut  $\mathcal{C} = \{S, T\}$ , and let  $e$  be an edge of positive capacity  $c_e$  crossing this cut from  $S$  to  $T$ . Decreasing the capacity of this edge by any  $\varepsilon > 0$  decreases the capacity of  $\mathcal{C}$  by  $\varepsilon$ . Since  $\mathcal{C}$  was a min-cut in the original graph, the min-cut in the new graph therefore has a strictly smaller min-cut and hence a strictly smaller max-flow (again, by weak duality). Thus, any positive capacity edge which crosses an  $s - t$  min-cut is a *critical edge*.

To find such an edge, we first compute the  $s - t$  max flow  $F$  of value  $f$  in  $G = (V, E)$ , and then construct the residual graph  $G'$ . If  $f = 0$ , then the network is disconnected (no path from  $s$  to  $t$ ) hence there are no-critical edges, so we report this and exit. Otherwise, let  $S$  be the set of vertices reachable from  $s$  in  $G'$  (we know that there is no path from  $s$  to  $t$  in  $G'$ , so  $S \neq V$ ). Since there are no edges from  $S$  to  $V - S$  in  $G'$ , it follows that  $f$  is equal to the capacity of the cut  $(S, V - S)$  in  $G$ . Thus, by the max-flow min-cut theorem,  $(S, V - S)$  is a minimum cut, and from our preceding discussion, we can simply return a positive capacity edge in  $G$  that goes from  $S$  to  $V - S$  (there exists such an edge since the capacity of the cut is  $f > 0$ ).

**Running time.** We first do a max-flow computation. After this, constructing  $G'$  takes  $O(|E| + |V|)$  time, and so does finding  $S$  (using BFS). Looking for an edge crossing  $(S, V - S)$  can take a further  $O(|V| + |E|)$  time, so the total running time is 1 max-flow computation +  $O(|V| + |E|)$  (for example,  $O(|V||E|^2)$  using Edmonds-Karp).

The proof of correctness is given below:

*Proof.* Let  $G = (V, E)$  be a flow network and let  $f^*$  be a maximum flow found using the Ford-Fulkerson algorithm. Let  $S_1$  and  $T_2$  be as defined above. We prove that the set of edges appear in every minimum  $(S, T)$ -cut is exactly

$$E_{\text{cut}} = \{(u, v) \in E \mid u \in S_1, v \in T_2\}.$$

First note that  $(S_1, V - S_1)$  is the minimum cut associated with the maximum flow. Furthermore,  $(V - T_2, T_2)$  is also a minimum cut because by the definition of  $T_2$ , all the edges from  $V - T_2$  to  $T_2$  must be fully saturated by  $f^*$ , and therefore  $v(f^*) = c(V - T_2, T_2)$ . The intersection of edges in  $(S_1, V - S_1)$  and edges in  $(V - T_2, T_2)$  is precisely  $E_{\text{cut}}$ . So  $E_{\text{cut}}$  must include all edges satisfying the requirement: Any edge not in  $E_{\text{cut}}$  is either not in the min-cut  $(S_1, V - S_1)$  or not in the min-cut  $(V - T_2, T_2)$ .

Now we only need to show that each edge  $(u, v) \in E_{\text{cut}}$  is indeed in all min-cuts. Consider an arbitrary min-cut  $(S, T)$ . Since we know  $s \in S$ , and there is a path from  $s$  to  $u$  in  $G_{f^*}$  (by definition of  $S_1$ ), if  $u$  is in  $T$ , then some edge along that path must be in the cut  $(S, T)$ . But this edge is not fully saturated by  $f^*$  (otherwise we wouldn't be able to find this path in  $G_{f^*}$ ), contradicting the assumption that  $(S, T)$  is a min-cut. Therefore,  $u$  must be in  $S$ . Similarly, we can show that  $v$  must be in  $T$ . Hence, the edge  $(u, v)$  is in the min-cut  $(S, T)$ .  $\square$

- 4. ( pts.) Application of Network Flow.** Some programs in our university have a complicated set of graduation rules. Suppose a set  $C = \{c_1, c_2, \dots, c_n\}$  of  $n$  courses is offered. Each rule is specified by a subset  $S \subseteq C$  of courses and an integer  $x$ , indicating that a student must take at least  $x$  courses from  $S$  in order to satisfy that rule. The subsets for different rules may overlap, but each course can be used to satisfy at most one rule. To graduate, a student must satisfy **all** of the rules.

Given a collection of  $m$  rules  $\{(S_1, x_1), (S_2, x_2), \dots, (S_m, x_m)\}$  and a set of  $\ell$  courses a student has taken  $\{c_{i_1}, c_{i_2}, \dots, c_{i_\ell}\}$ , describe an efficient algorithm to determine whether the student can graduate. Justify the correctness of your algorithm and analyze its running time.

For example, suppose there are 5 courses,  $C = \{c_1, c_2, c_3, c_4, c_5\}$ , and two rules:

- One must take at least  $x_1 = 2$  courses from  $S_1 = \{c_1, c_2, c_3\}$ .
- One must take at least  $x_2 = 2$  courses from  $S_2 = \{c_3, c_4, c_5\}$ .

Then a student who took  $\{c_1, c_3, c_5\}$  cannot graduate, while a student who took  $\{c_1, c_2, c_3, c_4\}$  can.

**Answer:** We build a flow network  $G$  as follows to represent the problem:

- One node for each course  $c_i$  [note: also OK to only include courses the student has taken], one node for each rule  $(S_j, x_j)$ , a source node  $s$ , a sink node  $t$ .
- Connect  $s$  to each course  $c_i$  with capacity 1 if the student has taken the course, otherwise 0 [note: equivalently, we can ignore the edge].
- Connect each rule  $(S_j, x_j)$  to  $t$  with capacity  $x_j$ .
- For each course  $c_i$ , connect it to all the rules  $(S_j, x_j)$  where  $c_i \in S_j$ , with capacity 1.

We then run any maximum flow algorithm on  $G$  to obtain a flow  $f^*$ . If the value of the maximum flow,  $v(f^*)$ , equals the total requirement  $\sum_{j=1}^m x_j$ , we conclude that the student can graduate.

For correctness, we show that  $v(f^*) = \sum_{j=1}^m x_j$  if and only if all graduation rules can be satisfied:

- ( $\implies$ ) Suppose there exists a flow  $f^*$  with  $v(f^*) = \sum_{j=1}^m x_j$ . Then all edges from the rule nodes to  $t$  must be saturated. Hence, for each rule node  $(S_j, x_j)$ , the total inflow equals  $x_j$ . By construction, this means there are  $x_j$  courses sending one unit of flow each to  $(S_j, x_j)$ , i.e., these courses are used to satisfy that rule. Because each course node has capacity 1 on its outgoing edges, no course can contribute to more than one rule. Therefore, all rules are satisfied without overlap.
- ( $\impliedby$ ) Conversely, suppose the student can graduate. Then there exist disjoint subsets  $X_1, X_2, \dots, X_m$  of the courses the student has taken such that  $X_j \subseteq S_j$  and  $|X_j| = x_j$  for all  $j$ . We can use this to construct a flow  $f$  on  $G$ : for each  $c \in X_j$ , send one unit of flow along the path  $s \rightarrow c \rightarrow (S_j, x_j) \rightarrow t$ . Because the  $X_j$ 's are disjoint and satisfy the required cardinalities,  $f$  is a valid flow. Moreover, it saturates all edges into  $t$ , so its value is  $v(f) = \sum_{j=1}^m x_j$ , implying  $f$  is a maximum flow.

The flow network  $G$  has  $O(n+m)$  nodes and  $O(n+m+\sum_{j=1}^n x_j)$  edges, both linear in the input size. Thus, the running time is dominated by the maximum flow computation. Using the Edmonds-Karp algorithm, the running time is  $O(|V||E|^2)$ , which is  $O(N^3)$  if  $N$  denotes the input size. [Note: it is OK to just say the running time is the same as a maximum flow algorithm without considering the input size of this problem.]

# **Rubric:**

**Problem 1, ? pts**

?

**Problem 2, ? pts**

?

**Problem 3, ? pts**

?

**Problem 4, ? pts**

?

**Problem 5, ? pts**

?